

SATS User Guide

Version: SATS 1.0.10

SATS (Signature Analyzer for Targeted Sequencing) is an R package for mutational signature analysis in targeted sequencing data. The method models panel-specific mutation opportunity, so it can be used for de novo signature detection, mapping of de novo profiles to tumor mutational burden (TMB)-normalized reference signatures, signature refitting in individual tumors, and calculation of signature-attributed mutation burdens.

The accompanying manuscript applies SATS to 111,711 tumors from AACR Project GENIE to construct a real-world, panel-calibrated pan-cancer catalogue of targeted sequencing-derived mutational signatures. The package is intended for targeted-panel cohorts where whole-exome or whole-genome sequencing data are unavailable or not routinely generated.

Installation

The current source version in this repository is SATS v1.0.10. Install the current GitHub source with:

```
if (!requireNamespace("devtools", quietly = TRUE))
  install.packages("devtools")
devtools::install_github(
  "binzhulab/SATS",
  subdir = "source",
  upgrade = "never"
)
```

Alternatively, download SATS_1.0.10.tar.gz from the repository and install the source archive:

```
R CMD INSTALL SATS_1.0.10.tar.gz
```

The source archive can also be installed from within R:

```
install.packages("../SATS_1.0.10.tar.gz", repos = NULL, type = "source")
```

SATS was formerly available from CRAN. CRAN currently lists the package as archived, so the GitHub source installation above is the recommended installation route for the current version. After installation, load the package with:

```
library(SATS)

if (packageVersion("SATS") < "1.0.10" ||
    !"GenerateVMatrix" %in% getNamespaceExports("SATS")) {
  stop("This guide requires SATS >= 1.0.10. Reinstall SATS and restart R.")
}
```

The de novo signature discovery examples use `signer`, which should be installed separately if that step is run locally. SATS can still be used for preprocessing, signature mapping, refitting and burden calculation without running `signer`.

Required Inputs

SATS uses three main input matrices. The mutation catalogue matrix V has dimension $P \times N$, where rows are mutation contexts and columns are samples. For SBS analysis, $P = 96$. The panel-context matrix L has the same dimension as V and gives the number of mutation opportunities per million base pairs for each mutation context and sample. The

reference signature matrix W has dimension $P \times K$, where columns are TMB-normalized reference signatures.

The row order of V , L and W must match. `ValidateSATSInputs()` can be used before analysis to check object type, numeric validity, finite values, non-negative values, integer-like mutation counts in V , duplicate or missing identifiers, mutation-context alignment, sample-ID alignment and signature-name alignment. If compatible named axes are present in different orders, SATS reorders them internally; if identifiers are absent from one paired object, duplicated or inconsistent, SATS stops with an informative error. For SBS analyses, SATS supports the COSMIC-style 96-channel order and the `signeR` order. The `SBS_order` argument in `GeneratePanelSize()` controls mutation-type ordering only; it does not select the COSMIC reference-signature version. The reference-signature version used by `MappingSignature()` is controlled separately by `COSMICv`, with "v3.4" as the current default.

The main workflow below starts from matched V and L matrices. If those matrices are already available, users can proceed directly to de novo signature detection, mapping, activity estimation and burden calculation. If users instead start from lower-level MAF-like mutation records and panel-coordinate information, `GenerateVMatrix()` and `GenerateLMatrix()` can first be used to construct matched V and L matrices. If users start from simple single-sample Variant Call Format (VCF) and Browser Extensible Data (BED) files, `ReadVCFAsMutationRecord()` and `ReadBEDAsPanelInfo()` can first convert those files into SATS-compatible tables. These preprocessing scenarios are described after the main workflow.

Example Data

The package includes simulated data in the `SimData` object:

```
data(SimData, package = "SATS")
names(SimData)
```

`SimData$V` is a simulated 96×10027 SBS mutation catalogue matrix and `SimData$L` is the corresponding sample-level panel-context matrix. These two matrices are already matched and are used for the main executable workflow. `SimData$TrueW_TMB` and `SimData$TrueH` are the simulated TMB-normalized signature profiles and activity matrix used to generate `SimData$V`. `SimData$PanelEx` and `SimData$PatientInfo` provide separate panel-coordinate and sample-annotation examples for the preprocessing special case below.

Main Analysis Workflow

The main analysis workflow starts from matched V and L matrices and proceeds through de novo signature detection, reference-signature mapping, sample-level activity estimation and signature-burden calculation.

Prepare Matched V and L Matrices

The workflow begins by assigning the simulated mutation-count matrix and panel-context matrix to `V_mat` and `L_mat`, then validating and aligning their mutation-context rows and sample columns.

```
data(SimData, package = "SATS")

V_mat <- SimData$V
L_mat <- SimData$L
dim(V_mat)
dim(L_mat)

validated <- ValidateSATSInputs(V = V_mat, L = L_mat)
V_mat <- validated$V
L_mat <- validated$L
```

In this main workflow, `L_mat` is the generated panel-context matrix used as the opportunity matrix for `signer()` and as the panel-context matrix for `EstimateSigActivity()` and `CalculateSignatureBurdens()`.

Detect De Novo Signatures

For cohort-level signature detection, SATS can be used with de novo profiles estimated by `signer` or another compatible signature extraction method. After `V_mat` and `L_mat` have been generated and aligned, choose the initial discovery strategy according to cohort size. If the sample size is small, for example fewer than 100 samples, the individual matched samples can be used directly. If the cohort is very large, such as a real-world cohort with about 10,000 tumors, every 100 matched samples can be pooled into one profile before de novo discovery. The executable workflow below uses the pooled strategy, because it mirrors the large-cohort setting used by the SATS manuscript. The individual-sample strategy is shown only as an optional commented block and is not required for the remaining examples. SATS stores mutation contexts in rows and samples in columns, whereas `signer()` expects samples in rows and mutation contexts in columns; therefore, the matched SATS matrices are transposed when passed to `signer()`. In the pooled workflow, `V_sum` and `L_sum` are derived by summing the same sample columns of `V_mat` and `L_mat`; they are not separate input files. In a full analysis, `W_hat` is the de novo TMB-normalized signature profile matrix returned by the extraction step. The examples below use simulated package matrices so that the code can be run end to end.

```
library(signer)

# Optional small-cohort strategy. This is commented out and is not run in the
# main guide workflow. Use it only when the cohort is small enough to fit
# individual samples directly.
# n_example <- min(100L, ncol(V_mat))
# sample_idx <- seq_len(n_example)
# V_fit <- V_mat[, sample_idx, drop = FALSE]
# L_fit <- L_mat[, sample_idx, drop = FALSE]
# stopifnot(identical(rownames(V_fit), rownames(L_fit)))
# stopifnot(identical(colnames(V_fit), colnames(L_fit)))
# set.seed(1)
# signer_re_100 <- signer(M = t(V_fit), Oport = t(L_fit), nlim = c(1, 5))
# W_hat_100 <- signer_re_100$Phat
# stopifnot(identical(rownames(W_hat_100), rownames(V_fit)))

# Pooled strategy: pool every 100 matched samples for a very large cohort.
# A cohort with about 10,000 tumors would yield about 100 pooled profiles.
pool_id <- ceiling(seq_len(ncol(V_mat)) / 100)
V_sum <- t(rowsum(t(V_mat), group = pool_id, reorder = FALSE))
L_sum <- t(rowsum(t(L_mat), group = pool_id, reorder = FALSE))
stopifnot(identical(rownames(V_sum), rownames(L_sum)))
stopifnot(identical(colnames(V_sum), colnames(L_sum)))

set.seed(1)
signer_re_pool <- signer(M = t(V_sum), Oport = t(L_sum), nlim = c(1, 5))
W_hat <- signer_re_pool$Phat
stopifnot(identical(rownames(W_hat), rownames(V_sum)))
```

Map De Novo Profiles to Reference Signatures

The de novo TMB-based profiles are then mapped to TMB-normalized reference signatures using `MappingSignature()`:

```
data(RefTMB, package = "SATS")

MappedSig <- MappingSignature(
  W_hat = W_hat,
  W_ref = RefTMB$TMB_SBS_v3.4
)
MappedSig
```

If `W_ref` is not supplied, `MappingSignature()` defaults to `COSMICv = "v3.4"` and uses `RefTMB$TMB_SBS_v3.4`. The returned data frame reports the selected reference signatures and the frequency with which each signature is selected across repeated penalized non-negative least-squares fits.

Estimate Signature Activities

After mapping the de novo profiles, use the mapped reference signatures for refitting. In this example, `SBS.list` is taken directly from `MappedSig$Reference`, so the activity and burden calculations are linked to the pooled `signer()` discovery and mapping result rather than to a manually specified signature list.

```
data(RefTMB, package = "SATS")

SBS.list <- MappedSig$Reference
if (length(SBS.list) == 0L)
  stop("No mapped signatures were selected; inspect W_hat or adjust mapping thresholds.")
if (!all(SBS.list %in% colnames(RefTMB$TMB_SBS_v3.4)))
  stop("At least one mapped signature is absent from the reference matrix.")

W_star <- as.matrix(RefTMB$TMB_SBS_v3.4[, SBS.list, drop = FALSE])

validated <- ValidateSATSInputs(V = V_mat, L = L_mat, W = W_star)
V_mat <- validated$V
L_mat <- validated$L
W_star <- validated$W

H_hat <- EstimateSigActivity(
  V = V_mat,
  L = L_mat,
  W = W_star
)
H_hat$H
```

`EstimateSigActivity()` uses an expectation-maximization algorithm and returns a list containing the estimated activity matrix H , the log-likelihood and a convergence flag. The estimated activity matrix has dimension $K \times N$.

Calculate Signature Burdens

Signature burdens are the expected numbers of mutations attributed to each selected signature in each tumor. They are calculated with `CalculateSignatureBurdens()`:

```
SigBdn <- CalculateSignatureBurdens(
  L = L_mat,
  W = W_star,
  H = H_hat$H
)

round(SigBdn[, 1:5], 2)
```

The returned matrix has dimension $K \times N$, with signatures in rows and samples in columns.

Special Case 1: Starting from MAF-like Mutation Records and Panel Information

If matched V and L matrices are not yet available, SATS provides preprocessing functions to construct them from lower-level inputs. `GenerateVMatrix()` generates an SBS96 or DBS78 mutation-count matrix from a MAF-like mutation record table. The mutation record must contain `Chromosome`, `Start_Position`, `End_Position`, `Variant_Type`, `Reference_Allele`, `Tumor_Seq_Allele2` and `Tumor_Sample_Barcode`. SBS analyses use records with `Variant_Type == "SNP"` and DBS analyses use records with `Variant_Type == "DNP"`.

```

dir <- system.file("extdata", "refitting_examples", package = "SATS")
sbs_file <- file.path(dir, "SBS_MAF_two_samples.txt")
sbs_mut <- read.table(sbs_file, header = TRUE, sep = "\t", quote = "",
                     stringsAsFactors = FALSE)

V_sbs <- GenerateVMatrix(sbs_mut, Class = "SBS", ref.genome = "hg19")
dim(V_sbs)

```

`GeneratePanelSize()` calculates panel-level mutation-context opportunity counts from panel-coordinate information. The input panel-coordinate data frame must contain Chromosome, Start_Position, End_Position and SEQ_ASSAY_ID. `GenerateLMatrix()` then converts panel-level context counts into a sample-level L matrix by matching sample identifiers to sequencing panels. The clinical sample table must contain SEQ_ASSAY_ID and either SAMPLE_ID or PATIENT_ID.

```

data(SimData, package = "SATS")

Panel_context_example <- GeneratePanelSize(
  genomic_information = SimData$PanelEx,
  Class = "SBS",
  SBS_order = "COSMIC",
  ref.genome = "hg19"
)

clinical_sample <- data.frame(
  SAMPLE_ID = unique(sbs_mut$Tumor_Sample_Barcode),
  SEQ_ASSAY_ID = SimData$PatientInfo$SEQ_ASSAY_ID[1],
  stringsAsFactors = FALSE
)

L_sbs <- GenerateLMatrix(
  Panel_context = Panel_context_example,
  Patient_Info = clinical_sample
)

if (!setequal(colnames(V_sbs), colnames(L_sbs)))
  stop("V and L contain different sample IDs")
if (!setequal(rownames(V_sbs), rownames(L_sbs)))
  stop("V and L contain different mutation-context rows")

L_sbs <- L_sbs[rownames(V_sbs), colnames(V_sbs), drop = FALSE]
identical(rownames(V_sbs), rownames(L_sbs))
identical(colnames(V_sbs), colnames(L_sbs))

```

The resulting `V_sbs` and `L_sbs` matrices can be used in the same main workflow shown above once sufficient samples are available for cohort-level signature discovery and mapping.

Optional VCF and BED Input Preparation

SATS v1.0.10 adds lightweight converters for users who start from standard targeted-panel input files. `ReadVCFAsMutationRecord()` converts a simple single-sample VCF file into the MAF-like mutation-record table used by `GenerateVMatrix()`. `ReadBEDAsPanelInfo()` converts BED target regions into the SATS panel-coordinate format used by `GeneratePanelSize()` and `GenerateLMatrix()`.

BED files use 0-based, half-open coordinates, whereas SATS panel-coordinate tables use 1-based, inclusive coordinates. `ReadBEDAsPanelInfo()` performs this conversion automatically by returning `Start_Position = BED_start + 1` and `End_Position = BED_end`.

```

dir <- system.file("extdata", "refitting_examples", package = "SATS")
vcf_file <- file.path(dir, "SBS_two_variants.vcf")
bed_file <- file.path(dir, "SATS_example_panel.bed")

mutation_record <- ReadVCFAsMutationRecord(vcf_file)
panel_info <- ReadBEDAsPanelInfo(

```

```

    bed_file = bed_file,
    seq_assay_id = "SATS_EXAMPLE_PANEL",
    name_col = 4
)

clinical_sample <- data.frame(
  SAMPLE_ID = unique(mutation_record$Tumor_Sample_Barcode),
  SEQ_ASSAY_ID = "SATS_EXAMPLE_PANEL",
  stringsAsFactors = FALSE
)

V_from_vcf <- GenerateVMatrix(
  mutation_record = mutation_record,
  Class = "SBS",
  ref.genome = "hg19"
)

L_from_bed <- GenerateLMatrix(
  Panel_context = panel_info,
  Patient_Info = clinical_sample,
  Class = "SBS",
  ref.genome = "hg19"
)

stopifnot(identical(rownames(V_from_vcf), rownames(L_from_bed)))
stopifnot(identical(colnames(V_from_vcf), colnames(L_from_bed)))

```

These converters are intended for input preparation, not for full clinical variant normalization. Complex VCF normalization, multi-sample genotype parsing, tumor-normal genotype interpretation, phasing and representation of complex events should be handled upstream when needed. The resulting SATS-compatible mutation-record and panel-coordinate tables can then be used with the standard SATS workflow.

Special Case 2: Single-Tumor or Small-Cohort Refitting

SATS can also refit a prespecified signature set in a single tumor or a small cohort. In this setting, the reference set should be constrained to signatures that are relevant to the tumor type and detectable in targeted sequencing data. `RefTMB$SBS_refSigs` and `RefTMB$DBS_refSigs` provide cancer-specific SBS and DBS reference-signature lists.

```

data(SimData, package = "SATS")
data(RefTMB, package = "SATS")

SBS.list <- RefTMB$SBS_refSigs[
  RefTMB$SBS_refSigs$cancerType == "Skin Cancer or Melanoma",
  "COSMIC"
]
W_star <- as.matrix(RefTMB$TMB_SBS_v3.4[, SBS.list])

V1 <- SimData$SingleTumorEx[, "singleV", drop = FALSE]
L1 <- SimData$SingleTumorEx[, "singleL", drop = FALSE]
colnames(V1) <- colnames(L1) <- "single_tumor"

H_hat <- EstimateSigActivity(V = V1, L = L1, W = W_star)
SigBdn <- CalculateSignatureBurdens(L = L1, W = W_star, H = H_hat$H)

```

This workflow supports targeted-sequencing refitting when the set of signatures is known from a cancer-type-matched catalogue or from a prior cohort-level SATS analysis.

Tests and Containerized Workflows

Unit tests are provided in `source/tests/testthat/` for `CalculateSignatureBurdens()`, `EstimateSigActivity()`, `ValidateSATSInputs()`, `GeneratePanelSize()`, `GenerateVMatrix()`, `GenerateLMatrix()`, `ReadVCFAsMutationRecord()` and `ReadBEDAsPanelInfo()`. These tests include

malformed-input checks for negative values, non-numeric values, duplicated identifiers, missing identifiers and mismatched named axes. To run them locally:

```
cd source
Rscript -e 'testthat::test_dir("tests/testthat")'
```

A Dockerfile is provided for building an R environment with SATS and its core dependencies installed:

```
docker build -t sats:1.0.10 .
```

The container can be used for a simple SATS smoke test:

```
docker run --rm sats:1.0.10 Rscript -e 'library(SATS); data(SimData, package="SATS"); stopifnot(is.matrix(SimData$V))'
```

The repository includes a shared `nextflow.config` with a Docker profile. The minimal Nextflow example in `nextflow/example1/` calls `GeneratePanelSize()` on an input `.rda` file containing `genomic_information` and optional `Class`, `SBS_order` and `ref.genome` objects:

```
nextflow run nextflow/example1/main.nf -profile docker \
  --outdir nextflow_results/example1 \
  --outfile outfile.rda
```

The expanded workflow in `nextflow/sats_workflow/` runs a compact SATS workflow from bundled package data. It validates matched V, L, W and H matrices with `ValidateSATInputs()`, maps example profiles with `MappingSignature()`, estimates signature activities with `EstimateSigActivity()` and calculates signature burdens with `CalculateSignatureBurdens()`:

```
nextflow run nextflow/sats_workflow/main.nf -profile docker \
  --outdir nextflow_results/sats_workflow \
  --n_samples 25
```

Expected outputs include `sats_mapping_results.csv`, `sats_activity_matrix.csv`, `sats_signature_burdens.csv`, `sats_workflow_outputs.rds` and `sats_workflow_summary.txt`. The example is intended as a reproducible workflow template. It does not run full de novo signature extraction and is not presented as a full clinical production pipeline or nf-core-compliant pipeline.

Citation and Web Resources

If you use SATS or the targeted-sequencing mutational-signature catalogue, please cite:

Lee et al., "A real-world pan-cancer catalogue of mutational signatures from 111,711 tumors" (submitted).

Software citation metadata are available in `CITATION.cff`.

Interactive catalogue plots are available at <https://analysistools.cancer.gov/mutational-signatures/##catalog/STS>. An online signature refitting tool is available at <https://analysistools.cancer.gov/mutational-signatures/##refitting>.